

Lecture #3: Metaheuristics

COMP3608 - Intelligent Systems

Inzamam Rahaman

Last Lecture

- Last lecture we looked at gradient descent for solving optimisation problems
- and backtracking for solving CSPs
- This lecture, we look at nature-inspired methods for solving such problems
- In particular, we shall look at 0th order methods that do not require the gradient

The Problem about problems

- Many optimisation problems are very difficult to solve!
- Sometimes, we have to acknowledge that we can't find optimal design point(s) to certain problems.
- Nevertheless, we can still try to get something passable or good, even if its not the best
 - Something is better than nothing

Heuristics

- Recall that optimisation problems are regarded as search problems
- Have a space of potential solutions
- Want to find design point that gives us optimal



Heuristics

- Sometimes, our problem has a structure that allow us to make good guesses about where good regions or good design points are located in the search space
- A rule of thumb for a very specific type of problem is called a heuristic
 - Example: When we have different servers at different costs for different tasks, we have heuristics that we can use to assign servers to tasks

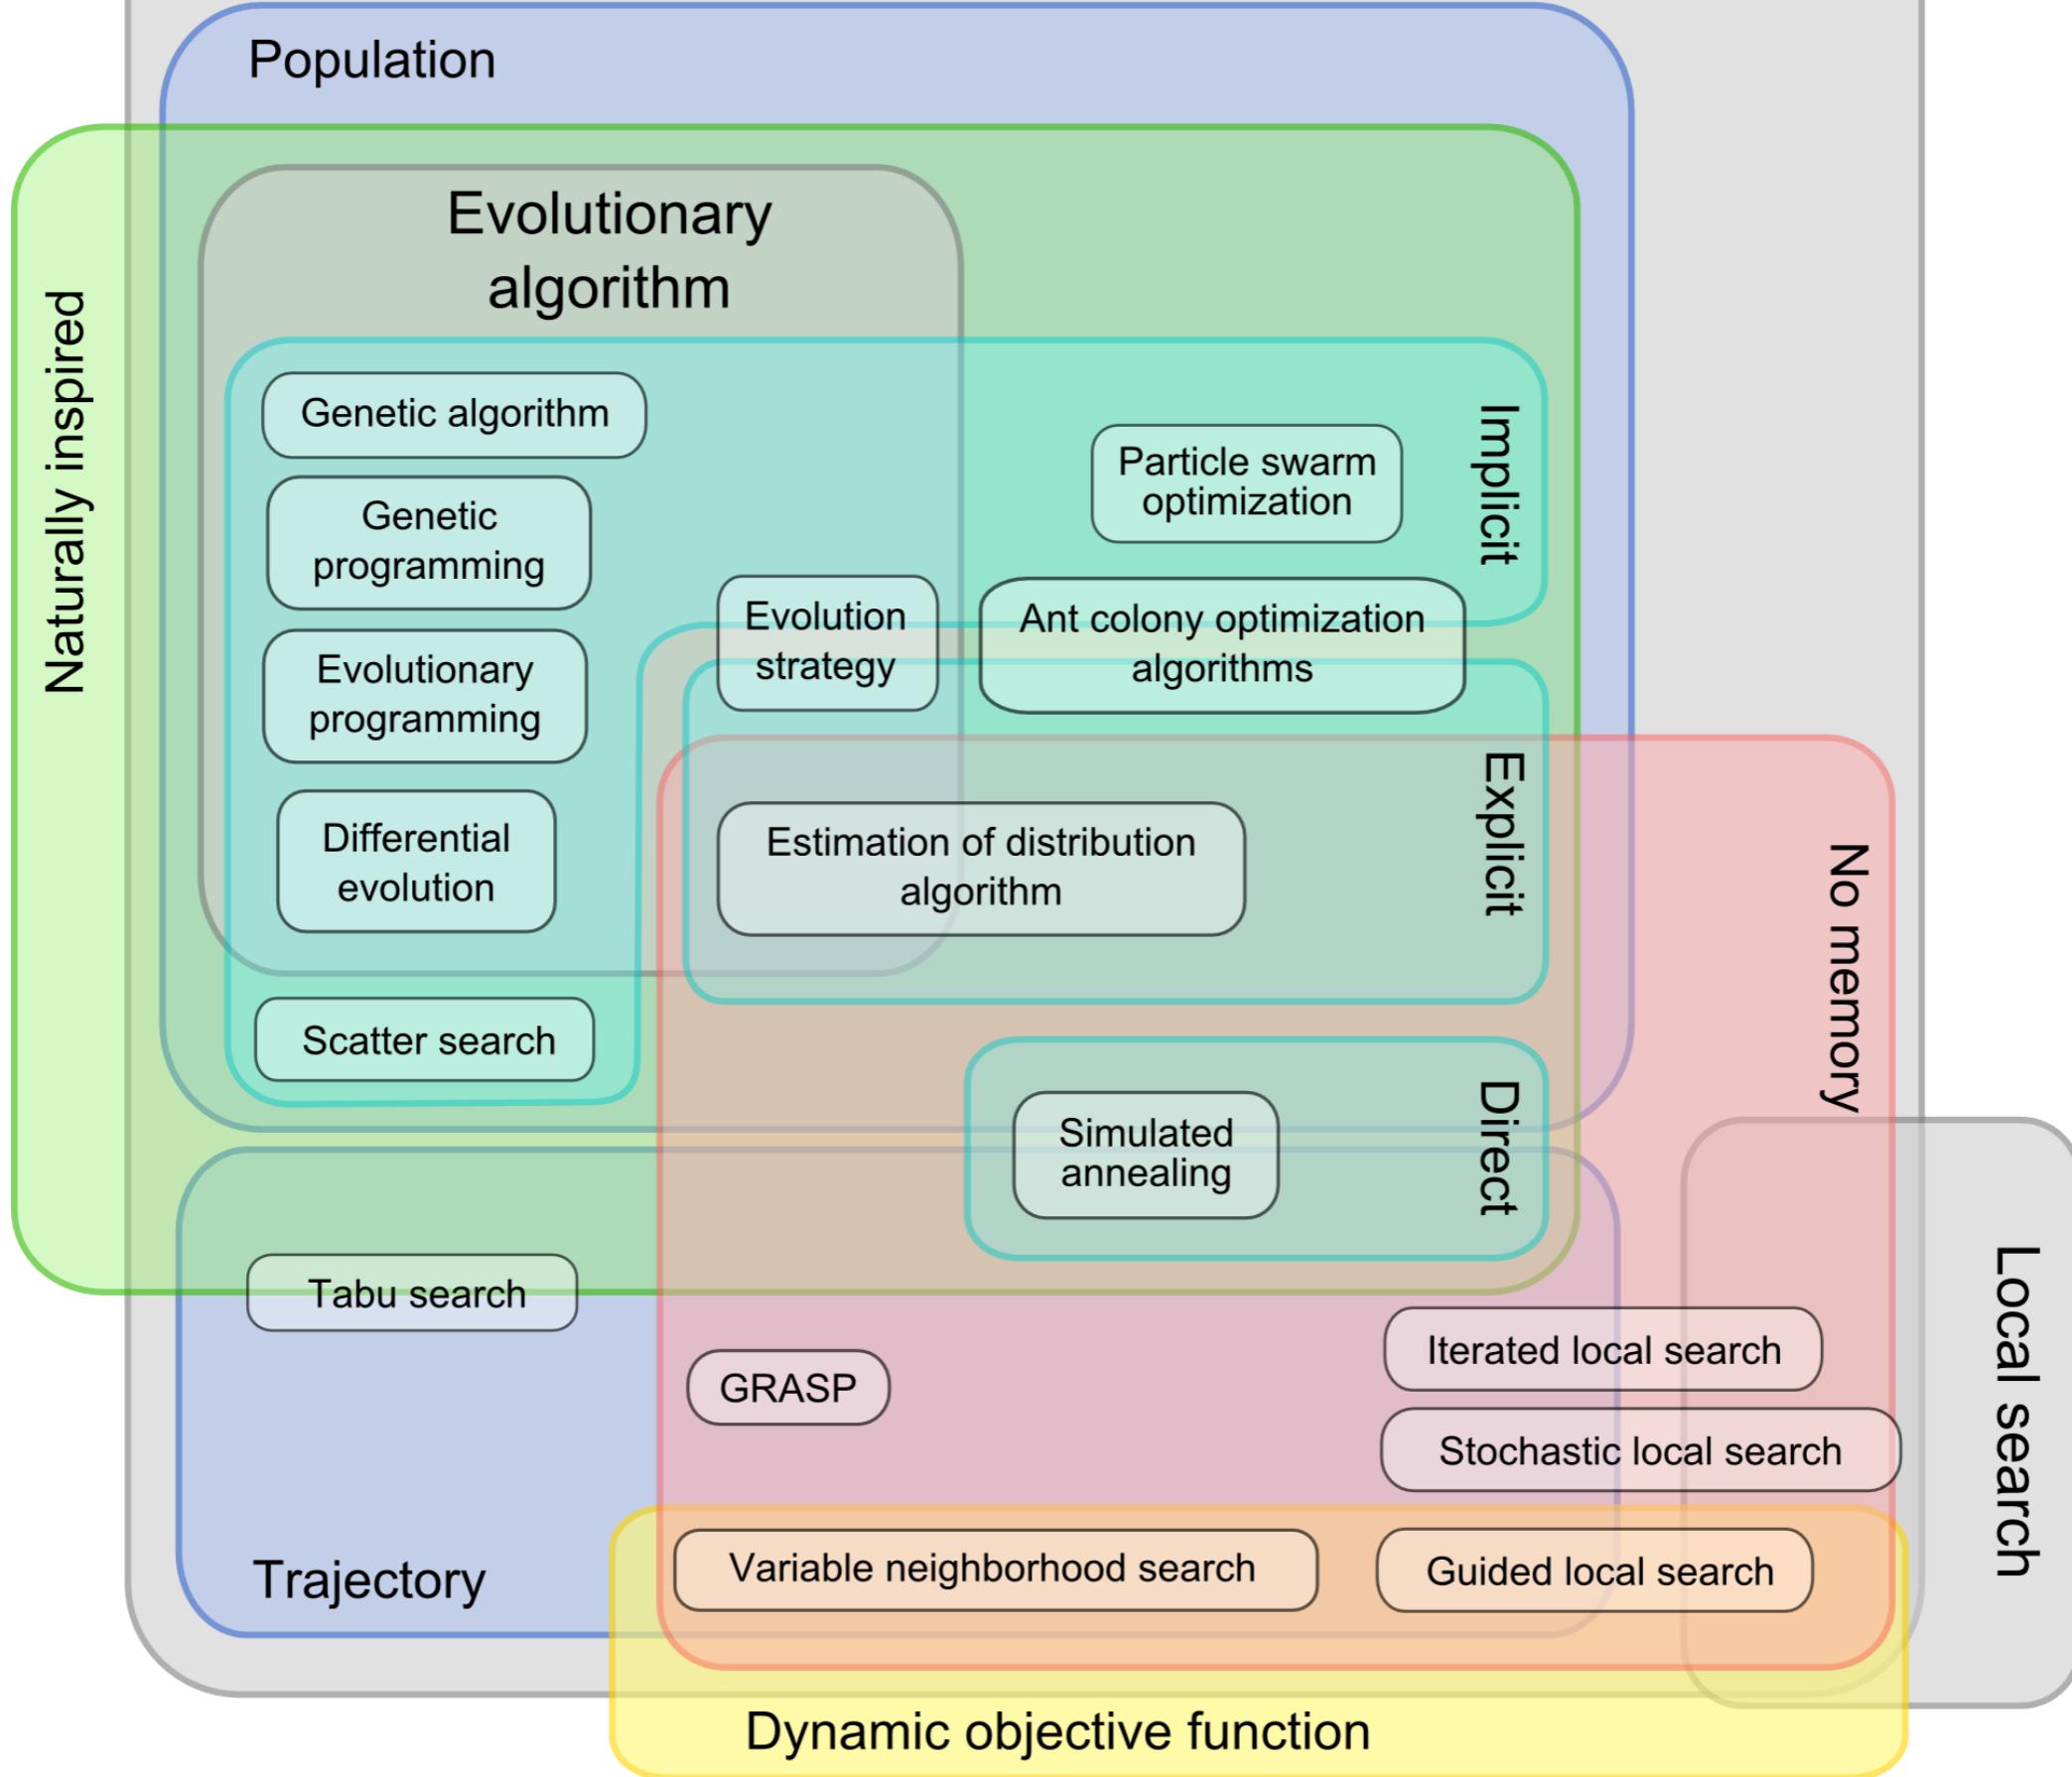
The problem with heuristics

- Heuristics, however, have problems
- For one, there are very problem and instance specific
 - Much effort must be applied to derive them
 - Much effort must be applied in their engineering
- Remember we like algorithms to be as general as possible!

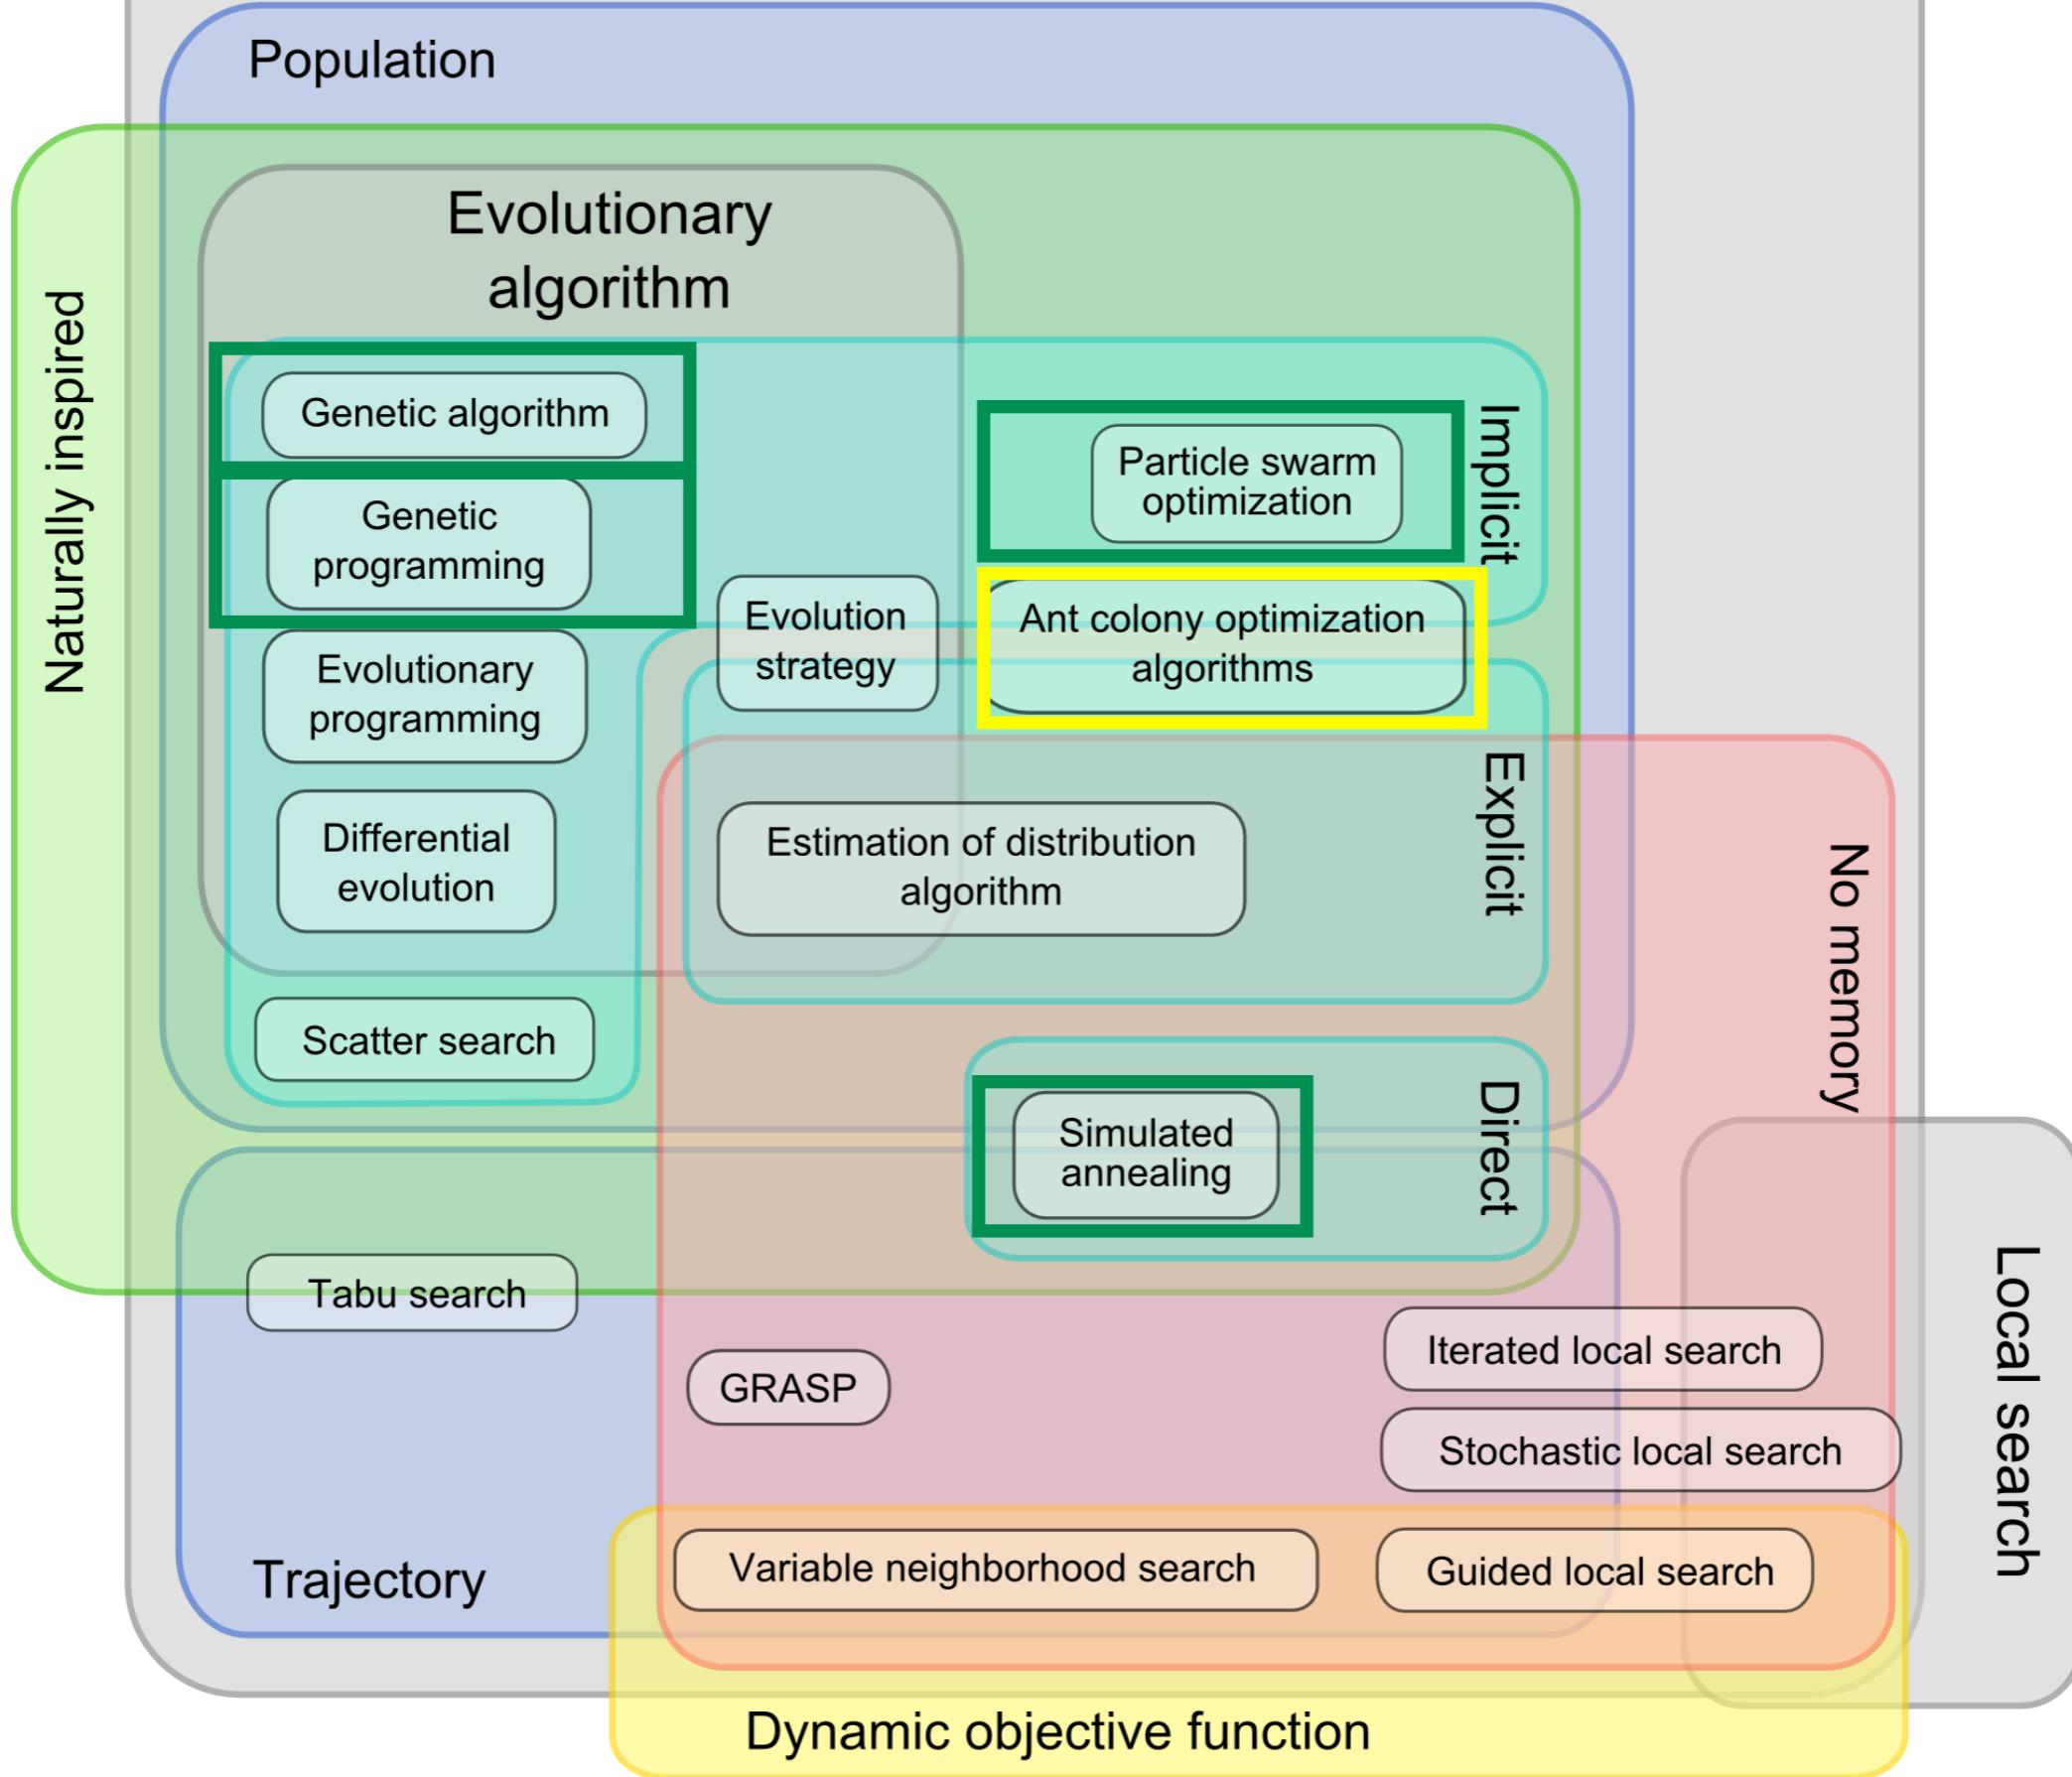
Metaheuristics to the rescue!

- Metaheuristics are like heuristics in that they offer good, efficient strategies but often no guarantees of optimality
- But they differ in that they are more generally applicable
- Rely on more generic ideas to guide search through space of potential solutions
- No magic bullet, but very helpful
- Many different types and general ideas
 - Exists an interesting taxonomy over metaheuristics

Metaheuristics



Metaheuristics



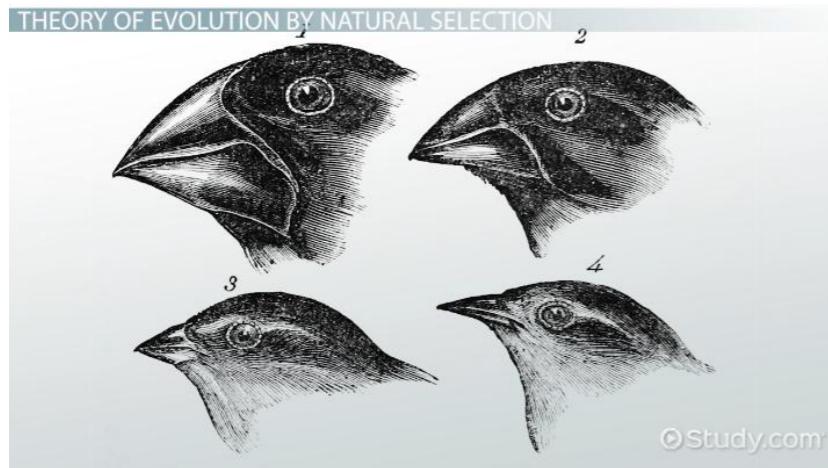
Metaheuristics

- Can have an entire course devoted to metaheuristics
- There are actually entire journals and conferences devoted just to metaheuristics
- GECCO is the most prominent venue in the area



The power of metaphor

- As you might have already guessed, many metaheuristics are nature-inspired
- Assume that the natural world solves optimisation problems (implicitly) all the time!



The power of metaphor

- Genetic Algorithm and Genetic Programming - Natural Selection
- Particle Swarm Optimisation - flocking behaviour of birds and insects
- Simulated Annealing - the cooling of molten metal into solid metal
- Ant Colony Optimisation - how ants forage for food and communicated indirectly by depositing pheromones

Genetic Algorithm

17-year-old crashes car while driving with a beanie pulled over her eyes as part of viral 'Bird Box' challenge

Alexandra Ma Jan 12, 2019, 7:19 AM



Natural Selection - A very brief summary

- Mechanism by which organisms evolve and develop traits
- Environment acts as a form of pressure on a population of organisms
- Organism's fitness is a result of how well it manages in its environment
- Organisms that survive mate
 - Fit organisms tend to want to mate with other fit organisms
 - Produce offspring
 - Offspring might acquire some mutation during gestation
- Offspring then continue the cycle



Natural Selection - A very brief summary

- Mechanism by which organisms evolve and develop traits
- Environment acts as a form of pressure on a **population** of **organisms**
- Organism's **fitness** is a result of how well it manages in its environment
- Organisms that survive mate
 - Fit organisms tend to want to **mate** with other fit organisms
 - **Produce** offspring
 - Offspring might acquire some **mutation** during gestation
- Offspring then **continue** the cycle

A bit of genetics

- Recall that the features of an organism are partially encoded by their genetics
- Hence, their fitness for their environment is at least partially determined by their genetics!
- DNA and chromosomes comprises of a sequence genes that work in concert with one another to encode features



If DNA/chromosomes is a sequence of genes,
from a CS perspective, what is DNA?



ARRAYS!

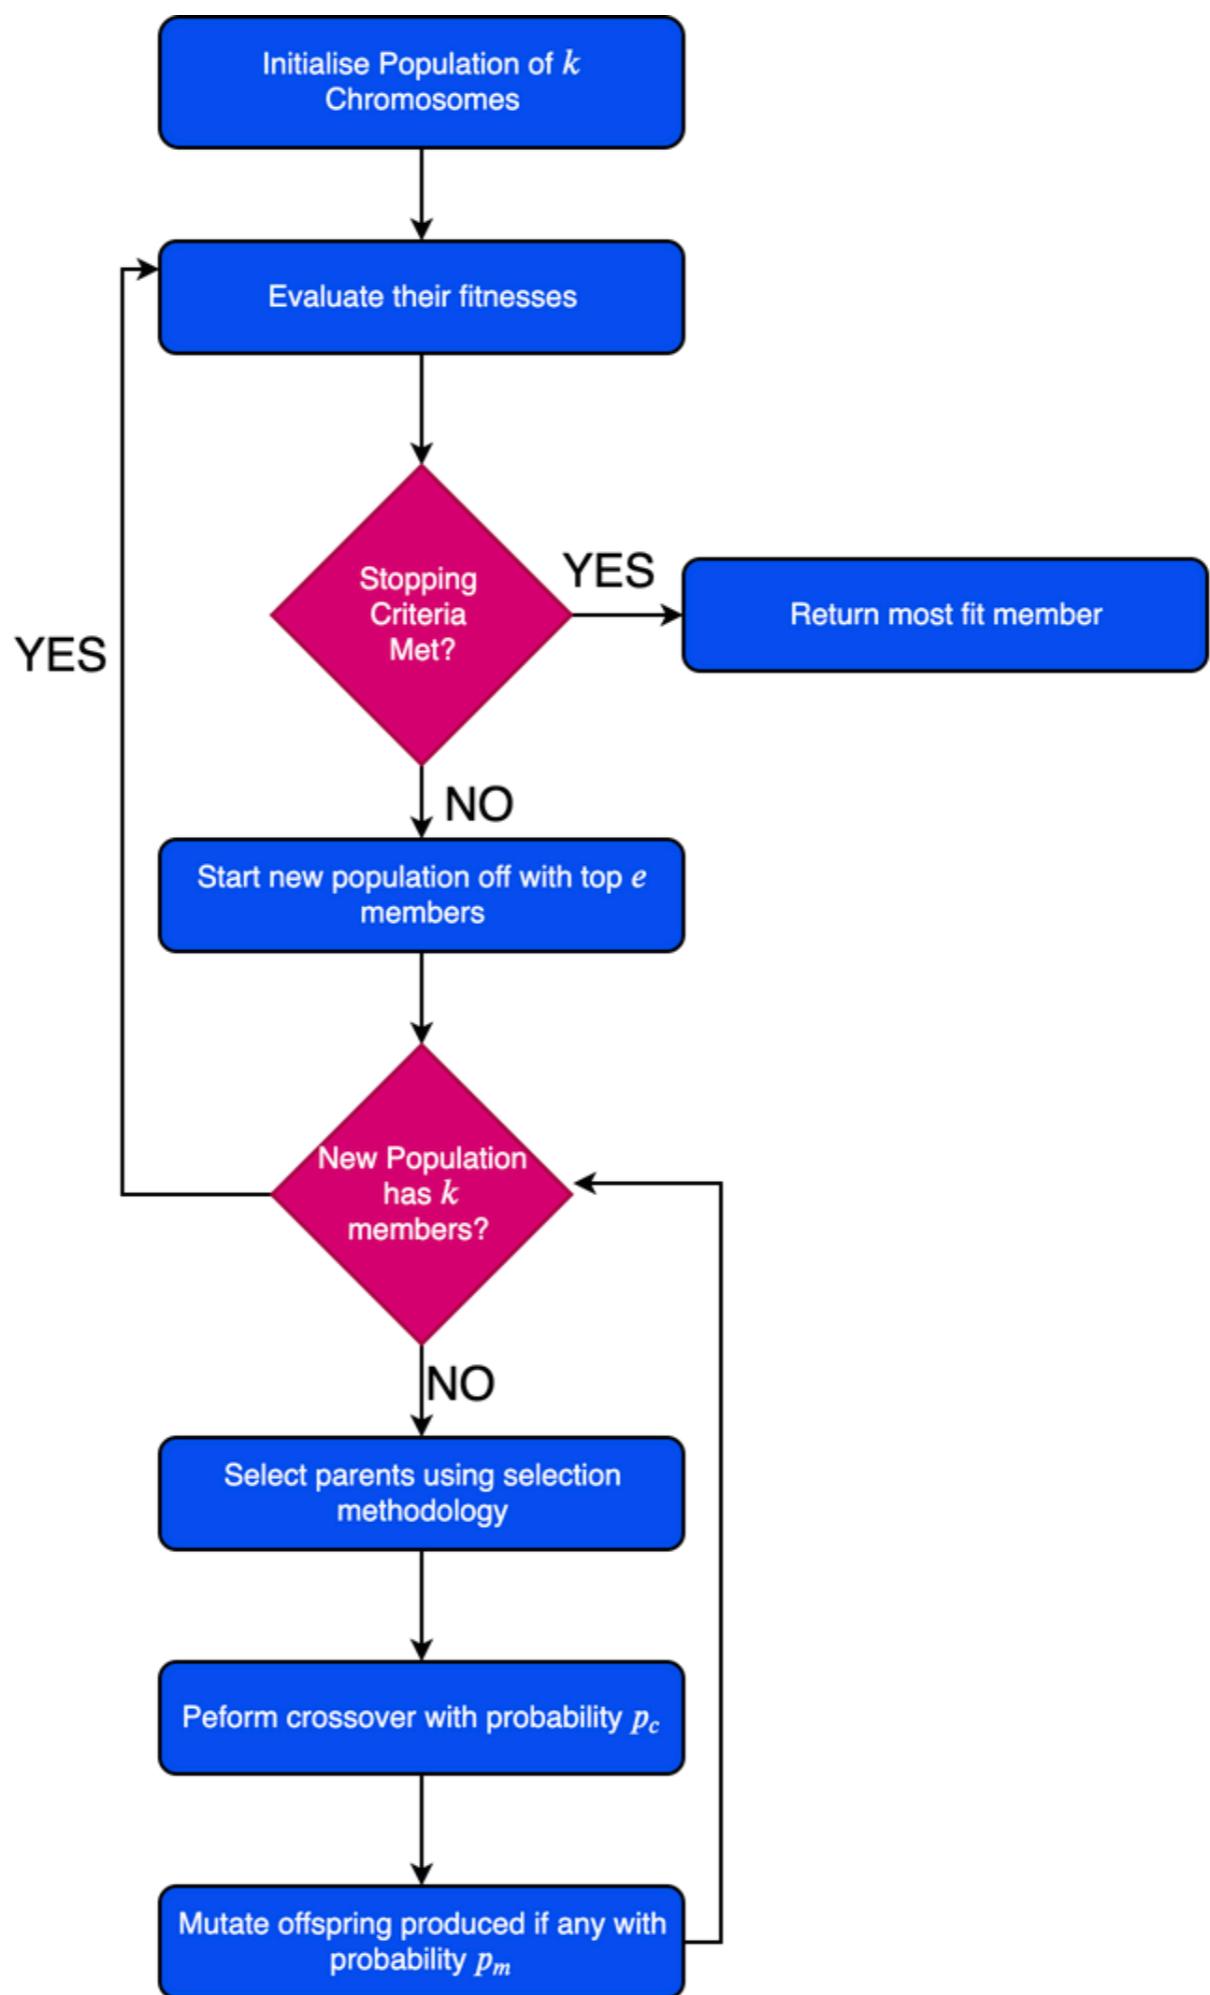


Key insights of genetic algorithm

- We can represent features (such as the components) of our design points as genes
- Represent the entire design point as an array of said features (like a chromosome)
- Use chromosome as input into some fitness function that assess how good the design point is

Key insights of genetic algorithm

- When organisms mate, their genetics are melded together - so we can combine our artificial chromosomes together similarly
- When organisms mutate, their genes change - can change values inside of the array
- We repeatedly apply this procedure to maximise some objective function that acts as our fitness function



Decisions, decisions, decisions

- The above is an outline
- We need to decide on some problem specific details
- Need to define
 - Representation
 - Crossover operation
 - Mutation
 - And of course fitness

Representation of Chromosomes

- We need to design an encoding of our design points as arrays.
- Each problem is different
 - Hence would have to select our representation to suite
 - Choosing a good representation is vital
 - Our choice of representation also dictates how we will randomly initialise our population
 - And how we apply other operations

Representation of Chromosomes

- Three main ways:
 - Real vectors - used for problems over real vector space. Each gene is component of vector is a real number representing a component
 - Binary arrays - each entry is 0 or 1 representing true or false for some feature associated with that entry. These can also be used to represent integer values
 - Permutations - the entire chromosome is a permutation

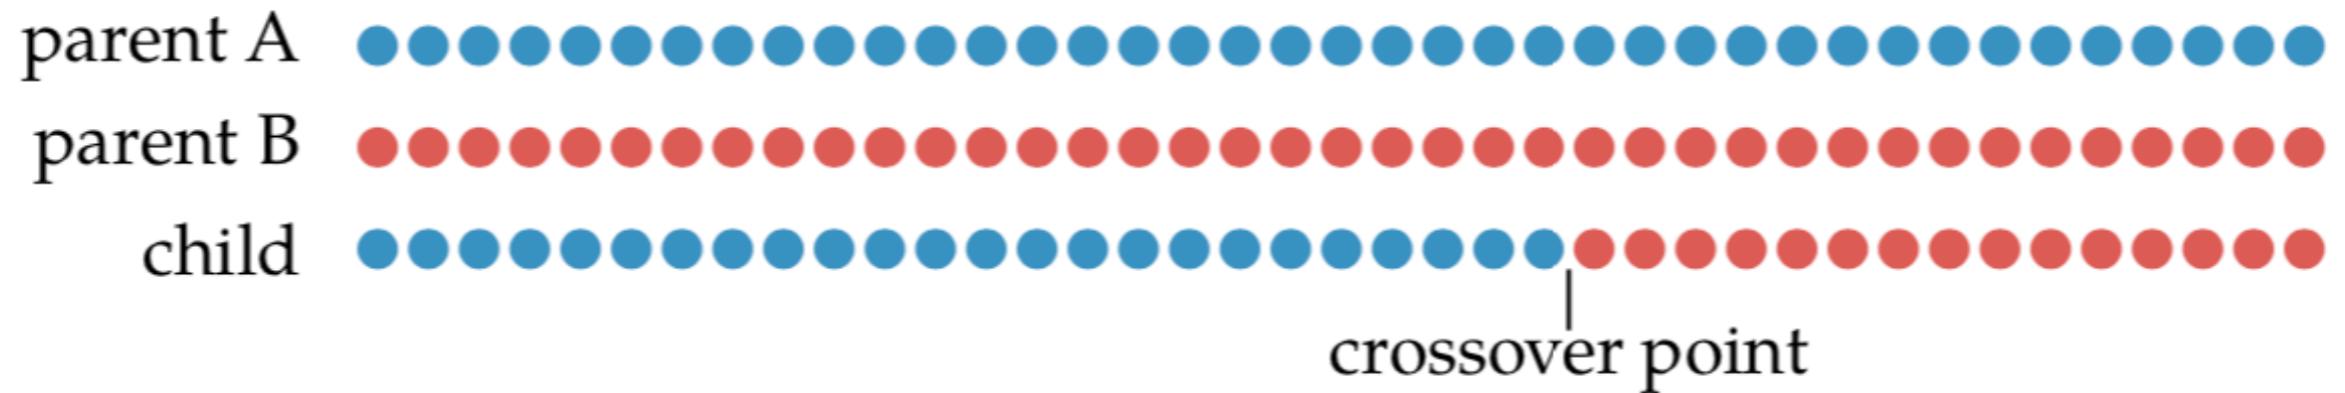
Selection strategy

- Roulette Wheel selection - all valid chromosomes in a population are considered for mating. Parents are selected stochastically in proportion to their fitness value
- Tournament Selection - t chromosomes are randomly selected. The top 2 are taken and then mated
- Truncation Selection - pick top t chromosomes, randomly choose 2 and mate them

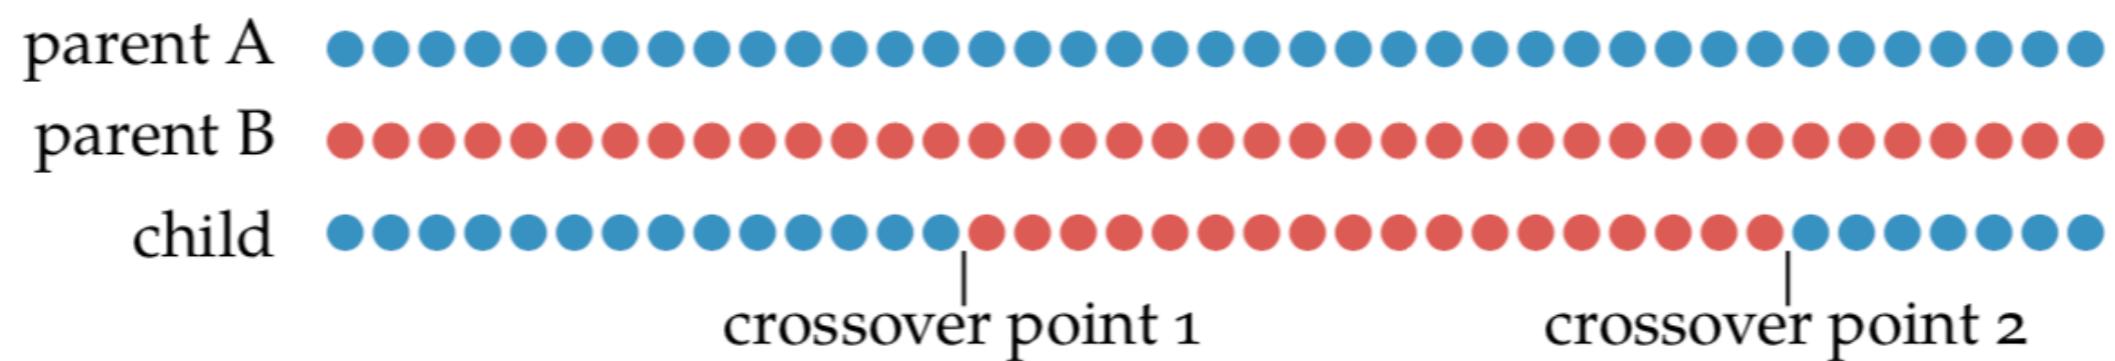
Crossover Methods

- Different methods of crossover
 - Single-point crossover
 - Multiple-point crossover
 - Uniform crossover
 - Interpolation crossover

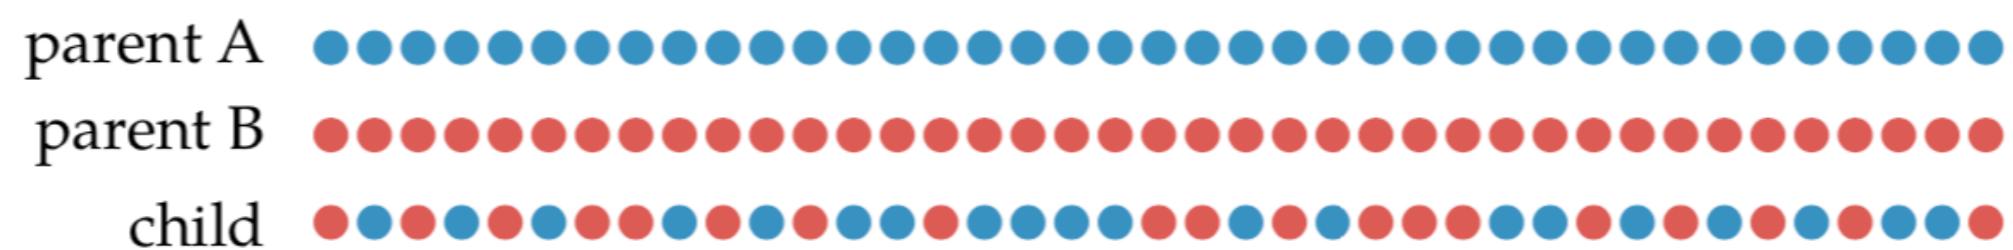
Single Point Crossover



Double-point crossover



Uniform Crossover

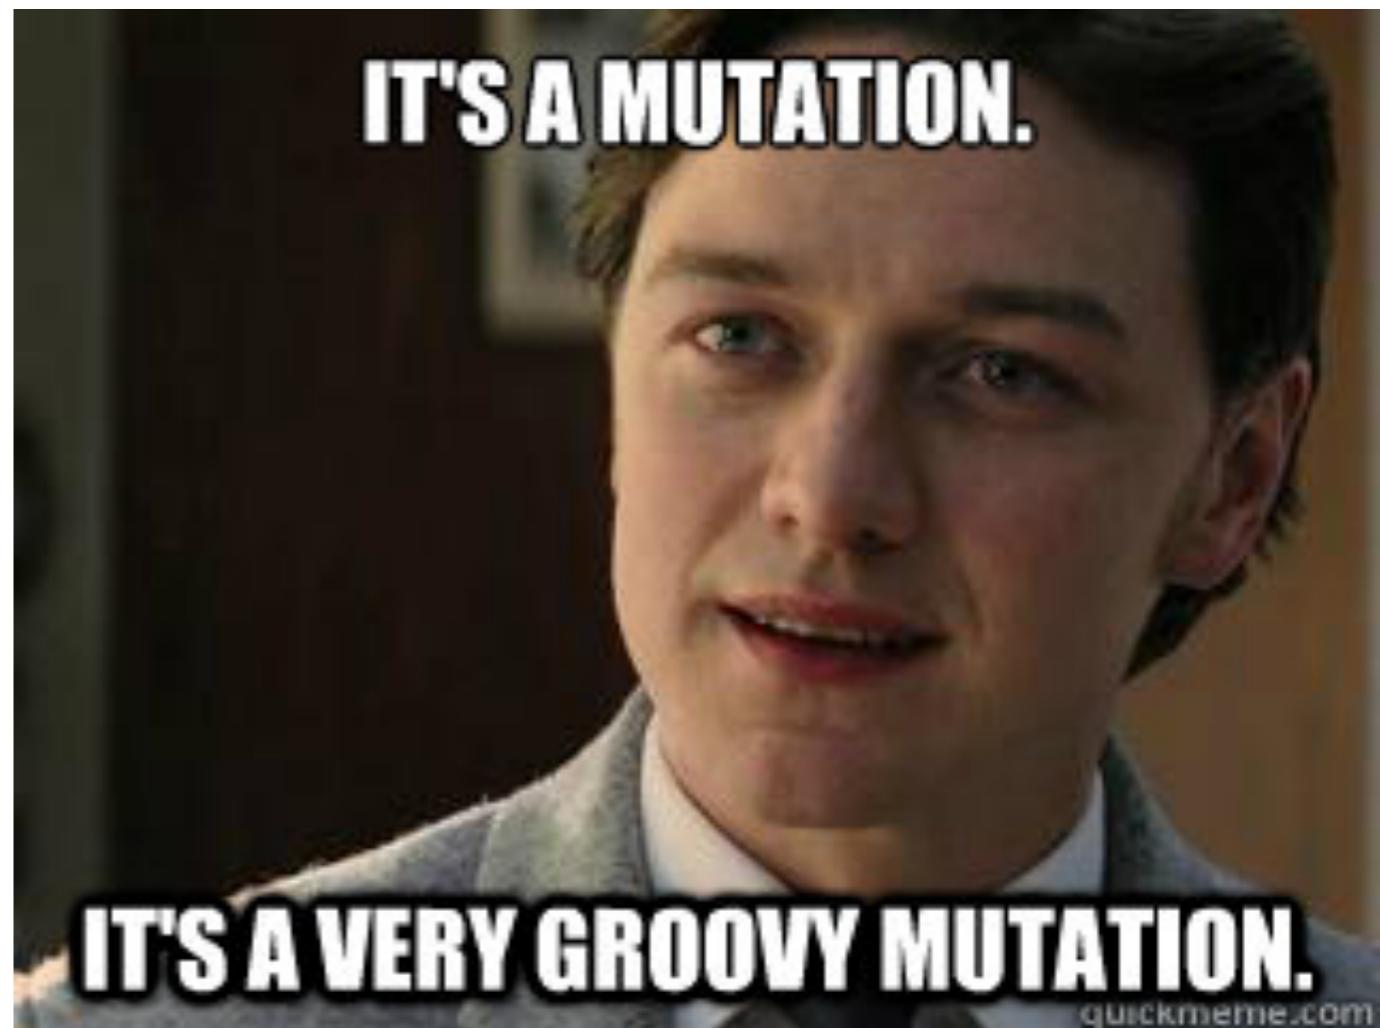


Interpolation Crossover

- Consider there being a line between two vectors x_1 and x_2 . This line is defined by the expression $(1 - \lambda)x_1 + \lambda x_2$ where $0 \leq \lambda \leq 1$
- We can look at crossover as choosing a point on this line.
 - We either randomly select a λ
 - Or set λ to a fixed value (usually 0.5)
- Interpolation crossover can only be done on chromosomes that are real vectors

Mutation

- Mutation can take on many different forms depending on the nature the representation
- Recall that our chromosome can be conceived as existing in some space of potential chromosomes
- Mutation makes this chromosome randomly jump to another neighbouring solution
- Mutation helps us avoid getting stuck in a bad local minimum

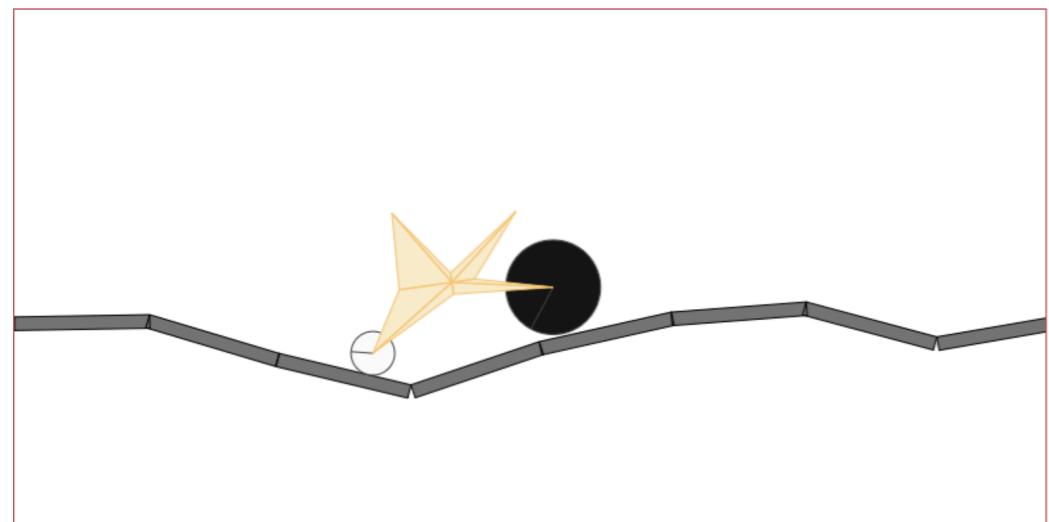


Mutation

- Possible interpretations
 - Adding standard Gaussian noise to a real vector chromosome
 - Flipping a bit in a bit-based chromosome
 - Swapping elements in a permutation based chromosome

GA Usage

- Using GAs require a deal of thought
- But they tend to work well when our search space is high, but number of feasible points is low
- Examples:
 - Scheduling
 - Automated design

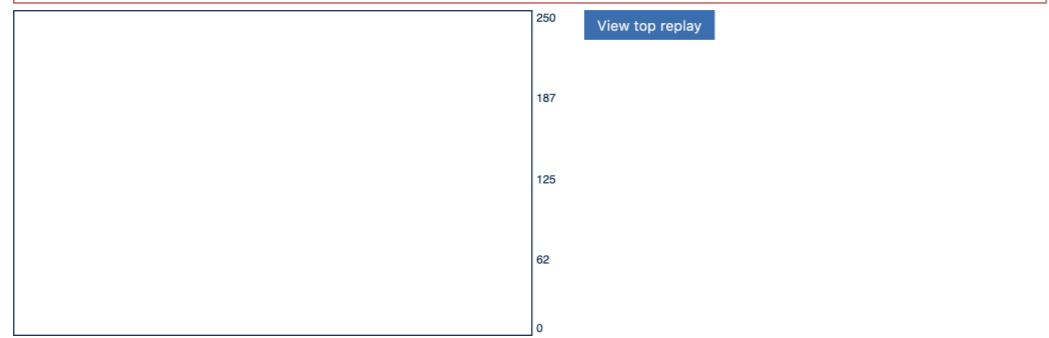


Donate

Save Population	Restore Saved Population
Surprise!	New Population
Fast Forward	

Create new world with seed
 Enter any string

Generation: 0
 Cars alive: 2
 Distance: 51.84 meters
 Height: -0.34 meters
 Mutation rate: 5%
 Mutation size: 100%
 Floor: fixed
 Gravity: Earth (9.81)
 Elite clones: 1



Watch Leader

†	
†	
2	—
†	
†	
6	—
†	
†	
†	
†	
†	
+	



Particle Swarm Optimisation

- Based off of flocking behaviour of birds and insects
- “Follow” the leader type strategy
- Each particle has a memory of its previous best state
- Each solution represented as a particle
 - Usually a real vector
- At the end of every iteration
 - Compute a velocity of each particle that gives it a direction to travel in
 - Add velocity to particle to compute new direction of particle
- Unlike GA, we can naturally use PSO for both maximisation and minimisation problems



Velocity Computation

- To decide on a particle's velocity, we take into account three pieces of information
 - Inertia - the resistance to change velocity
 - Exploration factor - tendency to explore the current globally best behaviour
 - Exploitation factor - tendency to follow previously known best personal behaviour

PSO Notation

- $x^{(i)}$ is the i th particle
- $v^{(i)}$ is the velocity of the i th particle
- $x_{\text{best}}^{(i)}$ is the best state of particle i in its entire history
- x_{best} is the state of the best performing particle over all iterations
- w - inertial factor, usually set to a value between 0 and 1
- r_1 - exploitation factor, randomly choose each time from $\mathcal{U}(0,1)$
- r_2 - exploration factor, randomly choose each time from $\mathcal{U}(0,1)$

$$\begin{aligned}v^{(i)} &= w v^{(i)} + r_1(x_{\text{best}}^{(i)} - x^{(i)}) + r_2(x_{\text{best}} - x^{(i)}) \\x^{(i)} &= v^{(i)} + x^{(i)}\end{aligned}$$

```
function pso(f, k, w, num_iter):
    population = []
    p_best = None
    best_score = -inf
    for i = 1 to k do
        population.append(init_particle())
    for i = 1 to num_iter:
        for p in population do
            f_p = f(p.state)
            if f_p > p.prev_best_score:
                p.prev_best = p.state
            if f_p > best_score:
                best_score = f_p
                p_best = p.state
        for p in population do:
            r1 = random_uniform(0, 1)
            r2 = random_uniform(0, 1)
            velocity = w * p.prev_velocity + r1(p.prev_best - p.state) + r2(p_best - p.state)
            p.state = p.state + velocity
            p.prev_velocity = velocity
    return p_best
```

Hill-Climbing Techniques

- GA and PSO are what are referred to as population based meta-heuristics
- They generate a random population, and use them in concert with one another to find a good design point
- Some methods, however, instead focus on randomly generating a single design point and iteratively refining it
- Hill-climbing techniques are a family of such techniques

Hill-Climbing Techniques

- Core idea, start at random point
- Each point has neighbouring points
- Examine some neighbouring points and test if they are better than our current point
 - If better move to that point
 - If equal or worse stay where we are
 - continue until converge criteria met



Hill-Climbing Techniques

- There is a major problem with vanilla hill-climbing techniques
- Vanilla hill-climbing techniques, more formally called deterministic hill climbing techniques can get stuck in a bad local optimal point
- Sometimes, we need to get worse before we get better
- Stochastic hill-climbing techniques avoid this problem by computing some probability of moving to a neighbouring point, and performing a Bernouilli Trial

Simulated Annealing

- SA is a variant of stochastic hill climbing
- Inspired by the process of smelting metal
- We provide a temperature schedule that controls volatility or willingness to accept a worse point in hopes of going to a better point in the future
- As time goes on (our iterations increase), temperature falls, thereby lowering volatility

Simulated Annealing - Probabilistic Jumping

- Suppose that we are at a point x_1 and we then jump to a point x_2
- Assuming that we want to maximise f , we let
$$\Delta f = f(x_2) - f(x_1).$$
- We go to point x_2 with the probability p where

$$p \begin{cases} 1 & \text{if } \Delta f \leq 0 \\ e^{-\frac{\Delta f}{t_i}} & \text{otherwise} \end{cases}$$

Simulated Annealing - Probabilistic Jumping

- Suppose that we are at a point x_1 and we then jump to a point x_2
- Assuming that we want to minimise f , we let $\Delta f = f(x_2) - f(x_1)$.
- We go to point x_2 with the probability p where

$$p \begin{cases} 1 & \text{if } \Delta f \leq 0 \\ e^{-\frac{\Delta f}{t_i}} & \text{otherwise} \end{cases}$$

← Metropolis Criteria

Simulated Annealing - Temperature Schedules

- Different types of temperature schedules
 - Fixed schedules - we come up with temperature values before hand
 - Non-fixed schedules - give starting temperature and parameters of schedule and temperature is dynamically calculated

Simulated Annealing - Temperature Schedules

- Exponential annealing:

- $t_i = \gamma^i t_0, \quad 0 < \gamma < 1$

- Fast annealing

- $t_i = \frac{t_0}{i}$

- Logarithmic annealing

- $t_i = \frac{t_0 \log(2)}{\log(i + 1)}$

```
function simulated_annealing(f, x, T, t, k_max)
    y = f(x)
    x_best, y_best = x, y
    for k in 1 : k_max
        x' = x + rand(T)
        y' = f(x')
        Δy = y' - y
        if Δy ≤ 0 || rand() < exp(-Δy/t(k))
            x, y = x', y'
        end
        if y' < y_best
            x_best, y_best = x', y'
        end
    end
    return x_best
end
```
